

CMPE 665 Homework 2

Simon Kirkwood & Adam Seidman
TA: Rohan Challa & Prangon Das
Lecture Section: 1
Professor: Muhammad Shaaban
Submission Date: December 2, 2021

By submitting this report, you attest that you neither have given nor have received any assistance (including writing, collecting data, plotting figures, tables or graphs, or using previous student reports as a reference), and you further acknowledge that giving or receiving such assistance will result in a failing grade for this course.

Your Signature: _____

Abstract

The goal of the assignment was to observe and analyze different partitioning methods for parallel programs via a ray tracing algorithm. Ray tracing is known for being highly parallelizable as each pixel of the image can be rendered entirely independently of the others, however the computation time of each pixel can vary wildly, so assigning pixels to separate processes in an efficient way is not a trivial task. Several different partitioning models were written and tested with the goal of comparing their performance and analyzing their strengths and weaknesses.

Design Methodology

For each partitioning model to which orientation applies, the horizontal orientation will be described, with the knowledge that the same principles apply in reverse to the vertical orientation. The first model written will be referred to as "static strips". It involves using the number of processes on which the program is run to divide the image into that many equal-sized rows or columns, depending on orientation.

First, the vertical resolution of the image is divided by the number of processes in order to calculate the height of each row. Then, each process uses its rank number as an index to determine its starting row and ending row, which is its rank plus one. This is multiplied by the row height to get the actual pixel values. The last process rank has its stop row set to the last row of pixels instead, in the event that the resolution was not evenly divisible by the number of processes. The processes then iterate through each of the rows to which they are assigned, shading every pixel in those rows.

The second algorithm is referred to as "static blocks." It is similar in principle to static strips in that it involves dividing the image into equal-sized chunks, however in this case it is divided into a grid of squares. It should be noted that for this reason, the amount of processes must have a square root that is a positive integer, as the grid dimensions are calculated using this square root.

Once the grid dimensions have been calculated (the width being equal to the height), the value is used to calculate the width and height in pixels of each block. From that, the rows and columns for each process are calculated. The starting row of a process is its rank

divided by the grid dimension, multiplied by the block height. The starting column is the remainder (modulo operation) of the rank divided by the block dimension, multiplied by the block width. The same iteration is then performed as with the static strips in order for each process to shade its assigned pixels.

The final static partitioning model is referred to as "static cycles." It revolves around splitting the image up into constant-sized rows/columns of which each process runs multiple. The process to which each strip is assigned follows a cyclical pattern so each process gets as close to an equal amount of strips as possible. It requires that the size of each cycle is defined as a command-line argument.

Using the predefined cycle size, each process uses its rank multiplied by the size as its starting row. Each time it loops, it adds this same value to itself in order to cycle to its next starting row. Within this loop, it goes row of pixels by row of pixels, shading each one across the entire width of the image. When the process has finished this for each cycle, it results in multiple thin separate strips of pixels going down the length of the image.

Every static model then has the same process for being combined. Each process send its shaded pixels to the master thread. The master thread then uses an identical algorithm to theirs, one by one, to determine for each process which pixels it was responsible for shading. It then takes each of those pixels and assigns its value to the master thread's own array of pixels. After doing this for every process in the communicator, the master thread has the complete image assembled, which it can then output to a file.

The dynamic partitioning model is designed to be slightly different. In this method, the user can supply command line arguments to determine the block sizes being sent to the worker processes. For this model, only the worker processes do pixel shading. The master process assigns blocks of pixels to the workers and when they are finished, it will assign them new ones. It keeps doing this until there are no blocks left to assign. Afterward, it waits until all other processes are finished and sends them a flag to tell them that they are done rendering the image. Each slave process then sends the pixels it rendered back to the master process, where they are reassembled into the main, final image.

Results/Analysis

As a control group, sequential run times were measured for both the simple and complex scene. For the simple scene, a time of 168.24 seconds was recorded. The complex scene yielded a time of 5799.96 seconds. Each partitioning model’s speedups are calculated in relation to these times.

The results of the first partitioning method, static strips, are shown in Table 1 and Table 2 for the simple and complex scenes respectively.

Table 1: Static vertical strips, rendering a simple scene.

Processes	Number of Strips	Execution Time	Speedup	C-to-C ratio
2	2	105.22	1.60	0.0040
4	4	67.44	2.49	0.0122
9	9	33.63	5.00	0.0490

The speedup for the simple scene is fairly strong with two processes, but the performance gain drops off as more processes are added. This is likely due to the scene having a variation in complexity for different parts of the scene, resulting in heavy load imbalance that puts much of the burden on relatively few processes.

Table 2: Static vertical strips, rendering a complex scene.

Processes	Number of Strips	Execution Time	Speedup	C-to-C ratio
2	2	5393.56	1.08	0.00011
4	4	3634.81	1.60	0.00014
9	9	2964.91	1.96	0.00007

The same holds true for the complex scene, this time to a very great extent. Running the scene with two processes has a speedup of only 1.08, indicating that one of the processes had much more work to do than the other, and this applies to an even larger extent with more processes added. Based on these results, statically splitting the image into strips is not well suited to ray tracing.

The next method is static blocks, represented in Table 3 and Table 4. For the complex scene, it should be noted that the cluster seemed to be stalling out during the time that it was attempted to be run with four processes, possibly due to a high volume of users at the time or some other issue, so only the results for nine processes were obtained.

Table 3: Static blocks rendering a simple scene

Number of Processes	Number of Blocks	Execution Time	Speedup	C-to-C
4	4	75.27	2.23	0.0249
9	9	52.19	3.22	0.0214

The speedup for the simple scene was quite weak for this partitioning method as well, again likely due to a heavy load imbalance in the complexity of the scene.

Table 4: Static blocks rendering a complex scene

Number of Processes	Number of Blocks	Execution Time	Speedup	C-to-C
9	9	3095.73	1.87	0.00034

The same once again holds true for the complex scene, to an even greater degree. The stark contrast in complexity between the edges of the scene and the center exemplify why this would occur.

For the last of the static methods, Table 5, Table 6, Table 7 and Table 8 represent results from the static cycles model.

Table 5: Static horizontal cycles, rendering a simple scene.

Processes	Height of Strip	Execution Time	Speedup	C-to-C ratio
4	1	48.27	3.48	0.0181
4	5	48.09	3.49	0.0173
4	10	48.44	3.47	0.0178
4	20	48.45	3.47	0.0191
4	80	48.88	3.44	0.0181
4	320	50.11	3.35	0.0174
4	640	54.36	3.09	0.0158
4	1280	82.78	2.03	0.0108
9	1	22.89	7.34	0.0894
9	5	22.79	7.38	0.0862
9	10	22.72	7.4	0.085
9	20	22.41	7.5	0.0805
9	80	21.99	7.65	0.0859
9	320	29.34	5.73	0.0565
9	640	46.11	3.64	0.0194
9	1280	81.27	2.07	0.0141

The cyclical model marks a significant improvement in speedup, for a sufficiently small cycle size. Because of the nature of the cycles, it is better at distributing the complex parts of the image out more evenly, so each process has a more equal amount of work to accomplish. This causes less time to be wasted, and therefore the entire execution finishes more quickly. This effect is lessened with larger cycle sizes because the image is broken up into chunks which are more likely to be separate in terms of complexity.

It should also be noted that with a cycle size high enough, there are not enough cycles to even assign to all 9 processes, so some go entirely without a workload in those cases, dramatically influencing the speedup negatively.

Table 6: Static horizontal cycles, rendering a complex scene.

Processes	Height of Strip	Execution Time	Speedup	C-to-C ratio
4	1	1490.41	3.89	0.00057
4	5	1547	3.74	0.00053
4	10	1520.3	3.81	0.00033
4	20	1687.25	3.43	0.00018
4	80	1773.2	3.27	0.0005
4	320	3173.03	1.82	0.00027
4	640	3012.96	1.92	0.00029
4	1280	3599.4	1.61	0.00014
9	1	666.66	8.7	0.00251
9	5	687.04	8.44	0.00173
9	10	729.27	7.95	0.00085
9	20	779.82	7.43	0.00229
9	80	1130.73	5.12	0.00123
9	320	2612.4	2.22	0.00063
9	640	3014.53	1.92	0.0003
9	1280	3613.55	1.6	0.00031

The same analysis applies to the complex scene once again. Its more dramatic differences in complexity between different areas of the scene accentuate these effects even more. Figure 1 shows the comparison between cycle size and execution time for both four and nine processes. Once the cycle height reaches 640, it can be seen that more processes become virtually useless.




Figure 1: Comparison of cycle sizes with 4 and 9 processes.

The actual cause of this is somewhat vague. A cycle size of 640 should result in up to eight processes continuing to provide speedup, yet this strangely seems not to be the case. Of course, for a cycle size of 1280, the cause of a lack of speedup is quite evident.

Table 7: Parallelism of static cycles with constant cycle size

Processes	Cycle Size	Execution Time	Speedup	C-to-C
2	27	94.38	1.78	0.0042
4	27	48.01	3.50	0.0182
9	27	22.65	7.43	0.0611

This table and Table 8 below are meant to illustrate how well the partitioning model scales with additional processes. Higher amounts of processes could not be run due to an issue with the cluster, however the few examples that could be taken illustrate that it scales quite well with additional processes, with the number of processes and the speedup being fairly proportional.

Table 8: Parallelism of static cycles with constant cycle size rendering a complex scene

Processes	Cycle Size	Execution Time	Speedup	C-to-C
2	27	2953.61	1.96	0.00019
4	27	1594.17	3.64	0.00036
9	27	808.642	7.17	0.00052

This once again holds true for the complex scene.

Dynamic partitioning worked quite well. As discussed earlier, the scenes can have dramatically different amounts of complexity depending on what is in each particular pixel. As a result, having a model that dynamically assigns groups of pixels can be very useful for keeping the load imbalance small. Table 9 and Table 10 show the results from this model.

Unfortunately, the 1x1 block size actually seemed to result in a segmentation fault, possibly due to too much memory being needed to store references to all the work units. This is a design flaw that could not be fixed in time and unfortunately caused results for that particular configuration to be unobtainable.

Table 9: Dynamic partitioning rendering a simple scene

Processes	Block Size	Execution Time	Speedup	C-to-C
9	1x1	N/A	N/A	N/A
9	15x15	26.12	6.44	0.0118
9	25x25	26.42	6.37	0.0114
9	50x50	26.17	6.43	0.0120
9	75x75	26.14	6.44	0.0113
9	100x100	26.04	6.46	0.0117

As expected, the dynamic model is great at reducing load imbalance as thus provides very impressive results. However, since it could only be run with nine processes, and not sixteen or thirty-six due to aforementioned cluster issues, the master thread not doing any calculations was actually a significant hamper on the speedup, proportionally speaking. As such, the speedup is actually slightly lower than for cycles. However, it could be expected that with a higher process count this effect would become less and less significant and dynamic partitioning would become the victor in terms of very high scaling.

Table 10: Dynamic partitioning rendering a complex scene.

Processes	Block Size	Execution Time	Speedup	C-to-C
9	1x1	N/A	N/A	N/A
9	15x15	739.35	7.84	0.00041
9	25x25	760.22	7.63	0.00037
9	50x50	800.29	7.25	0.00031
9	75x75	791.92	7.32	0.00035
9	100x100	848.04	6.84	0.00034

As expected, a model that is the best at minimizing load imbalance would excel for a complex scene such as this one. The speedups are very strong, but dip slightly with larger block sizes which increase load imbalance somewhat.

Generally, the static cycles model seems to be best for lower process counts. With many more processors available, however, the dynamic partitioning would likely scale better and provide a better speedup as a result since its load balancing cannot be matched.

Conclusion

The results of the exercise were successful, as they illustrated the clear differences in efficiency and scalability between different partitioning models. Through implementing and testing four different models, the requirements of scaling a raytracing algorithm become fairly evident. It also showcases a very prominent real-world example of parallelization which helps put the learned concepts into a better perspective. The issue with the cluster made the scope of the results limited, but this did not stop the data from showing some clear and illustrative patterns.